




MAPREDUCE TSQR

AUSTIN BENSON
arbenson@berkeley.edu



Author

Author

Author

TSQR

$A = QR$, $Q^T Q = I$, R upper triangular
Many more rows than columns \rightarrow "tall and skinny" (TS)
Lots of embarrassingly parallel work


Two methods for computing R
"TSQR" algorithm by Demmel et. al [5] (slower, more stable)
Cholesky decomposition on $A^T A$ (faster, less stable)

Make these algorithms run fast in the cloud [1], [4], [6]
Benchmark performance using Apache Hadoop
How does numerical stability factor in?

Data Serialization

TypedBytes storage [2], [3]
Different data types in sequence file
String format yields great improvement over list format


Packed rows
Store 2 or 4 rows per record
Performance can be significantly better, stay the same, or get worse



Fault Tolerance

Key advantage to Hadoop and other MapReduce architectures
How does this affect performance?
How does $P(\text{fault})$ affect performance?


Noticeable but manageable
Faults quickly introduce a 25% performance hit
 $P(\text{fault}) \approx 1/5 \rightarrow$ only 50% performance hit!



Peak Performance

How close to peak performance?
1-3% peak for large matrices
For a MapReduce architecture, this is about what we expect

Refining the model
 ~ 60 seconds for launch, cleanup overhead
 ~ 60 MB/s disk reads (1 TB SATA disks)
 $\rightarrow \sim 7-11\%$ peak



Numerical Experiments



Text Text
Text
Text

Text Text
Text
Text

Strong Scaling

Difficult to control processor allocation
In Hadoop, use `mapred.min.split.size` parameter
Embarassingly parallel work greedily consumes any extra resources

For TSQR, computation time small
Hadoop overhead and disk reads dominate time



MapReduce schemes for TSQR (L) and Cholesky (R)

Streaming: C++ vs. Python

Use Hadoop streaming
Python provides easy prototyping for testing algorithms at a large scale (both algorithms implemented in 100 lines of code)
C++ can give us better performance

Text Text
Text
Text

Text Text
Text
Text

Future work

Implement customized data storage
Expand to other areas of linear algebra (e.g., LU)
Explicit formulation of Q
Experiment with other MapReduce frameworks (e.g., Spark, Twister)

References

- [1] Austin Benson. MapReduce TSQR code. <https://github.com/arbenson/mrtsqr>
- [2] Klaas Bosteele. Dumbo. <https://github.com/kibosteele/dumbo>
- [3] Klaas Bosteele. TypedBytes. <https://github.com/kibosteele/typedbytes>
- [4] Paul G. Constantine and David F. Gleich. *Tall and Skinny QR factorizations in MapReduce architectures. MAPREDUCE 2011.*
- [5] James Demmel, Laura Grigori, Mark F. Hoemmen, and Julien Langou. *Communication-optimal parallel and sequential QR and LU factorizations.* UCB/EECS-2008-89, August 2008.
- [6] David Gleich. MapReduce TSQR code. <https://github.com/dgleich/mrtsqr>