




MAPREDUCE TSQR

AUSTIN R. BENSON
arbenson@berkeley.edu

TSQR


- $A = QR$, $Q^T Q = I$, R upper triangular
 - Many more rows than columns \rightarrow "tall and skinny" (TS)
 - Lots of embarrassingly parallel work
- Two methods for computing R
 - "TSQR" algorithm by Demmel et al. [5] (slower, more stable)
 - Cholesky decomposition on $A^T A$ (faster, less stable)
 - Both algorithms scale well
- Make these algorithms run fast in the cloud [1], [4], [6]
 - Implementations with Apache Hadoop MapReduce
 - How does numerical stability factor in?



MapReduce schemes for TSQR (left) and Cholesky QR (right)


Data Serialization

- TypedBytes storage [2], [3]
 - Different data types in sequence file
 - String format yields great improvement over list format
- Packed rows
 - Store 2 or 4 rows per record
 - Performance can be significantly better, stay the same, or get worse




Strong Scaling

- Difficult to control processor allocation
 - In Hadoop, use `mapred.min.split.size` parameter
 - Embarrassingly parallel work greedily consumes any extra resources
- For TSQR, computation time small
 - Hadoop overhead and disk reads dominate time




Peak Performance

- How close to peak performance?
 - 1-3% peak for large matrices
 - For a MapReduce architecture, this is about what we expect
- Refining the model
 - ~60 seconds for launch, cleanup overhead
 - ~60 MB/s disk reads (1 TB SATA disks)
 - \rightarrow 7-11% peak




Fault Tolerance

- Key advantage to Hadoop and other MapReduce architectures
 - How does this affect performance?
 - How does P(fault) affect performance?
- Noticeable but manageable
 - Faults quickly introduce a 25% performance hit
 - $P(\text{fault}) \sim 1/5 \rightarrow$ only 50% performance hit




Numerical Experiments

- How stable are our algorithms?
 - Compute R with TSQR and Cholesky
 - $Q = AR^{-1}$, check $\|Q^T Q - I\|_2$
 - Q not quite orthogonal $\rightarrow A = QR = Q'R'R$



Streaming: C++ vs. Python

- Use Hadoop streaming
 - Python provides easy prototyping for testing algorithms at a large scale (both algorithms implemented in 100 lines of code)
 - C++ can give us better performance
- Number of columns matters
 - 10 columns \rightarrow about the same performance
 - 200 columns \rightarrow C++ runs 4x faster than Python



Future work

- Implement customized data storage
- Expand to other areas of linear algebra (e.g., LU) and tensor algebra
- Explicit formulation of Q
- Experiment with other MapReduce frameworks (e.g., Spark, Twister)

References

- [1] Austin Benson. MapReduce TSQR code. <https://github.com/arbenson/mrtsqr>
- [2] Klaas Bosteele. Dumbo. <https://github.com/kibosteele/dumbo>
- [3] Klaas Bosteele. TypedBytes. <https://github.com/kibosteele/typedbytes>
- [4] Paul G. Constantine and David F. Gleich. Tall and skinny QR factorizations in MapReduce architectures. In Proceedings of the second international workshop on MapReduce and its applications, MapReduce '11, pages 43-50. New York, NY, USA, 2011. ACM.
- [5] James Demmel, Laura Grigori, Mark F. Heemmen, and Julien Langou. Communication-optimal parallel and sequential QR and LU factorizations. UCB/ECS-2008-89, August 2008.
- [6] David Gleich. MapReduce TSQR code. <https://github.com/dgleich/mrtsqr>