



MAPREDUCE TSQR


AUSTIN R. BENSON
arbenson@berkeley.edu

TSQR

$A = QR$, $Q^T Q = I$, R upper triangular
Many more rows than columns \rightarrow "tall and skinny" (TS)
Lots of embarrassingly parallel work

Two methods for computing R
"TSQR" algorithm by Demmel et. al [5] (slower, more stable)
Cholesky decomposition on $A^T A$ (faster, less stable)
Both algorithms scale well

Make these algorithms run fast in the cloud [1], [4], [6]
Implementations with Apache Hadoop MapReduce
How does numerical stability factor in?




MapReduce schemes for TSQR (left) and Cholesky QR (right)

Data Serialization

TypedBytes storage [2], [3]
Different data types in sequence file
String format yields great improvement over list format


Packed rows
Store 2 or 4 rows per record
Performance can be significantly better, stay the same, or get worse



Strong Scaling

Difficult to control processor allocation
In Hadoop, use `mapred.min.split.size` parameter
Embarrassingly parallel work greedily consumes any extra resources


For TSQR, computation time small
Hadoop overhead and disk reads dominate time



Peak Performance

How close to peak performance?
1-3% peak for large matrices
For a MapReduce architecture, this is about what we expect

Refining the model
~60 seconds for launch, cleanup overhead
~60 MB/s disk reads (1 TB SATA disks)
 \rightarrow 7-11% peak




Fault Tolerance

Key advantage to Hadoop and other MapReduce architectures


How does this affect performance?
How does P(fault) affect performance?

Noticeable but manageable
Faults quickly introduce a 25% performance hit
 $P(\text{fault}) \sim 1/5 \rightarrow$ only 50% performance hit




Numerical Experiments

How stable are our algorithms?
Compute R with TSQR and Cholesky
 $Q = AR^{-1}$, check $IQ^T Q - I$
 Q not quite orthogonal $\rightarrow A = QR = Q'R'R$



Streaming: C++ vs. Python

Use Hadoop streaming
Python provides easy prototyping for testing algorithms at a large scale (both algorithms implemented in 100 lines of code)
C++ can give us better performance



Number of columns matters
10 columns \rightarrow about the same performance
200 columns \rightarrow C++ takes 4x faster than Python

Future work

Implement customized data storage
Expand to other areas of linear algebra (e.g., LU) and tensor algebra
Explicit formulation of Q
Experiment with other MapReduce frameworks (e.g., Spark, Twister)

References

- [1] Austin Benson. MapReduce TSQR code. <https://github.com/arbenson/mrtsqr>
- [2] Klaas Bosteele. Dumbo. <https://github.com/kibosteele/dumbo>
- [3] Klaas Bosteele. TypedBytes. <https://github.com/kibosteele/typedbytes>
- [4] Paul G. Constantine and David F. Gleich. Tall and skinny QR factorizations in MapReduce architectures. In Proceedings of the second international workshop on MapReduce and its applications, MapReduce '11, pages 43-50. New York, NY, USA, 2011. ACM.
- [5] James Demmel, Laura Grigori, Mark F. Heemmen, and Julien Langou. Communication-optimal parallel and sequential QR and LU factorizations. UCB/ECS-2008-89, August 2008.
- [6] David Gleich. MapReduce TSQR code. <https://github.com/dgleich/mrtsqr>